

# Initial Problem

Hot spots:

- Occurs any time when large number of clients wish to access the data from single server
- Multiple requests causes swamping, rendering it useless.

# Terminologies

$c$ : Set of caches

$C$ : Number of caches

$p$ : Set of all pages

$\delta(m_1, m_2)$ : Latency - time taken for a message from machine  $m_1$  to arrive at  $m_2$

$q$ : Number of requests a cache must see before bothered to store a copy of page

# Random Trees

- Simplifications to the model:
  1. All machines know about all caches
  2.  $\delta(m_i, m_j) = 1$  for all  $i \neq j$
  3. All requests are made at same time
- Static Model - Only one batch of requests
- Basic idea:
  - extension of “tree of caches” - use a different, randomly generated tree for each page
  - ensures that no machine is near the root for many pages- load balancing

# Random Trees

## Protocol :

- Associate a rooted d-ary tree (abstract tree) with each page
- Request form a 4-tuple: identity of the requester, name of desired page, sequence of nodes, sequence of caches
- Mapping :
  - root always mapped to server for page
  - other nodes are mapped to caches by hash function:

$h: p \times [1.....C] \rightarrow c$  - distributed to all browsers and caches

- Parameter q avoid copies of pages with fewer requests

# Random Trees

- Browser (wants a page)
  - picks a random leaf to root path
  - maps nodes to caches with  $h$
  - asks the leaf node for the page
  - request - name of browser, name of page, path, result of mapping
- Cache (receives a request)
  - first checks if it is caching a copy of page or is in the process of getting one to cache
  - if so, it returns the page to the requester
  - otherwise, it increments the counter for page and node it is acting as and asks the next cache on path
  - if counter reaches  $q$  - caches the copy of page
- Server (receives a request)
  - sends the requester a copy of page

# Random Trees

## Latency:

- If request is forwarded from leaf to root
  - latency =  $2 \log_d C$
- Latency is less if the request is satisfied with a cached copy
- In practice, time required to obtain large page is not multiplied by the number of steps in path it travels

# Random Trees

## Swamping :

- The following theorem provides high probability bounds on the number of requests a cache gets.

*if  $h$  is chosen uniformly and at random from the space of functions  $p \times [1 \dots C] \rightarrow c$ , then with probability at least  $1 - 1/N$ , where  $N$  is a parameter, the number of requests a given cache gets is no more than*

$$\rho(2\log_d C + O(\log N / \log \log N)) + O(dq \log N / (\log(dq/\rho * \log N)) + \log N)$$

# Random Trees

## Proof:

### 1. Request to leaf nodes

- Each request coming for page  $p$  goes to a random leaf node in that page's tree
- $L$  denotes the no. of leaf nodes in each abstract tree
- Associate a weight  $w_p = r_p / L$  with each abstract leaf node of  $p$ 's tree
- Machine  $m$  has  $1/C$  chance of being at an arbitrary leaf node
- Let  $V_{pj}$  = event when  $j^{\text{th}}$  leaf node is assigned to  $m$
- $V_{pj} = 1$ , with probability  $1/C$  and  $V_{pj} = 0$  otherwise
- Bounding the total weight  $W = \sum_p w_p V_{pj}$  assigned to  $m$



# Random Trees

- We would use Chernoff Bound here.
- $W$  is weighted sum of poisson variables with weights possibly greater than 1
- But, each weight  $w_p = r_p / L \leq p/d(d-1)$
- So we can apply chernoff bounds to  $W/p/d(d-1)$
- This gives bound of  $O(p \log N / \log \log N)$  on  $W$  which holds with probability  $1 - 1/N$

## 2. Request to internal nodes

- No internal node gets more than  $dq$  requests
- Since only  $R$  requests in all, bound the no. of abstract nodes that get  $dq$  requests

# Random Trees

- In fact bound,  $n_j$ , the no. of nodes over all trees that receive between  $2^j$  and  $2^{j+1}$  requests where  $0 \leq j \leq \log(dq)-1$
- $\sum r_p \leq R$
- Since, each of the  $R$  requests gives at most  $\log_d C$  requests, total no. of requests is no more than  $R \log_d C$
- So,  $\sum_{j=0}^{\log(dq)-1} 2^j n_j \leq R \log_d C$

Lemma: The total no. of internal nodes that receive at least  $qx$  requests is at most  $2R/x$ ,  $x > 1$

Proof:

- Any node gets at most  $q$  requests from each child, a node that gets at least  $qx$

# Random trees

requests must have downward degree of at least  $x > 1$ .

- Consider all nodes  $u$  with downward degree one
- Let  $v$  and  $w$  be parent and child of  $u$  respectively
- Replace all such nodes with degree 1 by a single edge from  $v$  to  $w$ .
- Since  $x > 1$ , we now have a tree where each node has downward degree of at least 2
- The no. of leaves in tree is no more than no. of requests,  $r_p$
- So, if there are  $y$  nodes with downward degrees of at least  $x$ , the  $xy \leq 2r_p$  and so  $y \leq 2r_p/x$
- Thus, total no. of nodes over all trees which receive at least  $qx$  requests is no more than  $\sum 2r_p/x = 2R/x$

# Random Trees

- For  $x=1$ , there can clearly be no more than  $R \log_d C$  requests
- The lemma showed:  $n_j$  is at most  $2R/2^j$ , but for  $j=0$ ,  $n_j$  is at most  $R \log_d C$
- Assignments of nodes are independent, the prob that a machine  $m$  gets more than  $z$  of these nodes is at most  $\binom{n_j}{z} (1/C)^z (en_j/C z)^z$
- In order for RHS to be as small as  $1/N$ , we set

$$z = \Omega(n_j/C + \log N / \log((C/n_j) \log N))$$

- Latter will be present only if  $C/n_j \log N > 2$ , so  $z$  is  $O(n_j/C + \log N / \log((C/n_j) \log N))$  with prob at least  $1 - 1/N$ .
- So, with prob at least  $1 - \log(dq)/N$ , the total no. of requests received by  $m$  due to internal nodes is of order of

# Random Trees

- $\sum_{j=0}^{\log(dq)-1} 2^{j+1} (n_j/C + \log N / \log((C/n_j) \log N))$

$$= 2 \rho \log_d C + O(dq \log N / \log((dq/\rho) \log N)) + \log N \text{ ----- } 2$$

hence from 1 and 2

$$\rho (2 \log_d C + O(\log N / \log \log N)) + O(dq \log N / \log((dq/\rho) \log N)) + \log N$$

# Consistent Hashing


Problem: Evenly distributing objects across caches such that each client should know where to query for a particular object.

Typical Solution: The server can use a hash function to evenly distribute objects across caches.

But,

- When a new cache is added or removed, every item would be hashed to a different location.

# Consistent Hashing




A particular client needs to only aware of a set of caches.

- In *consistent hashing* a hash function is viewed as a ring
- Each server has its own key region (its "position" on the ring)

For example

- The key region for 2nd server is the area between 1 and 2
- Only 2 is responsible for keys in that region.

# Consistent Hashing



Suppose we want to add a new server

- we just pick some area and divide it into 2 parts
- Assign the new server one of these two.
- The keys that happen to be in that region are moved to the new server
- Each server knows the key range which it manages
- We can route the request to the server that is close to the key we're looking for



# Consistent Hashing

## Virtual Nodes

There are some challenges with this basic approach

- random position assignment may lead to non-uniform data/load distribution
- heterogeneity in performance is assumed (that is, we assume that all the servers have same performance)


A variant of Consistent Hashing algorithm addresses this issue:

- instead of mapping a single node to the ring, each node gets multiple points there. (*virtual nodes*)

A virtual node looks like a single node, but it refers to the real node.

Advantages: Load balancing.

# Consistent Hashing



Properties of consistent hashing.

- "spread"
  - "load"
  - "smoothness"
  - "balance"
  - "monotonicity"
- 
- Ranged hash function:  $f : 2^\beta \times I \rightarrow \beta$
  - View: The set of caches of which a particular client is aware.  $V \subseteq \beta$
  - $f(v, i) = f_v(i)$  = the bucket to which item  $i$  is assigned in view  $v$ .

# Consistent Hashing

Balance: Distributing items among buckets in a balanced fashion.

A ranged hash family is balanced if for a given view ( $V$ ), a set of items, and a randomly chosen hash function, with high probability, the fraction of items mapped to each bucket is  $O(1/|V|)$

Monotonicity: Item may move from an old bucket to a new bucket , but not from one old bucket to another. (intuition about consistency)

A ranged hash function  $f$  is monotone if for all views  $V_1 \subseteq V_2 \subseteq \beta$ ,

$fv_2(i) \in V_1$  implies  $fv_1(i) = fv_2(i)$ .

# Consistent Hashing

Spread: Let  $V_1 \dots V_v$  be the set of views, altogether containing  $C$  distinct buckets and each individually containing at least  $C/t$  buckets. The spread  $\sigma(i)$  is the quantity  $|\{f_{v_j}(i)\}_{j=1}^v|$ . The spread of a hash function is the maximum spread of an item. When each person tries to assign an item  $i$  to a bucket, there are at most  $\sigma(i)$  different opinions about which bucket should contain this item. Good consistent hashing function have low spread.

Load: For a ranged hash function  $f$  and bucket  $b$ , the load  $\lambda(b)$  is the quantity

$|\bigcup_v f_v^{-1}(b)|$ . The load of a hash function is the maximum load of a bucket. There are at most  $\lambda(b)$  distinct items that at least one person thinks belongs to the bucket.

